

Centroid-Based Document Classification: Analysis & Experimental Results*

Eui-Hong (Sam) Han and George Karypis

University of Minnesota, Department of Computer Science / Army HPC Research Center
Minneapolis, MN 55455
{han, karypis}@cs.umn.edu

Abstract. In recent years we have seen a tremendous growth in the volume of text documents available on the Internet, digital libraries, news sources, and company-wide intranets. Automatic text categorization, which is the task of assigning text documents to pre-specified classes of documents, is an important task that can help both in organizing as well as in finding information on these huge resources. Text categorization presents unique challenges due to the large number of attributes present in the data set, large number of training samples, and attribute dependencies. In this paper we present a simple linear-time centroid-based document classification algorithm, that despite its simplicity and robust performance, has not been extensively studied and analyzed. Our experiments show that this centroid-based classifier consistently and substantially outperforms other algorithms such as Naive Bayesian, k -nearest-neighbors, and C4.5, on a wide range of datasets. Our analysis shows that the similarity measure used by the centroid-based scheme allows it to classify a new document based on how closely its behavior matches the behavior of the documents belonging to different classes. This matching allows it to dynamically adjust for classes with different densities and accounts for dependencies between the terms in the different classes. We believe that this feature is the reason why it consistently outperforms other classifiers that cannot take into account these density differences and dependencies.

* This work was supported by NSF CCR-9972519, by Army Research Office contract DA/DAAG55-98-1-0441, by the DOE ASCI program, and by Army High Performance Computing Research Center contract number DAAH04-95-C-0008. Access to computing facilities was provided by AHPARC, Minnesota Supercomputer Institute. Related papers are available via WWW at URL: <http://www.cs.umn.edu/~karypis>

1 Introduction

We have seen a tremendous growth in the volume of online text documents available on the Internet, digital libraries, news sources, and company-wide intranets. It has been forecasted that these documents (with other unstructured data) will become the predominant data type stored online. Automatic text categorization [22, 18, 13, 4, 9], which is the task of assigning text documents to pre-specified classes (topics or themes) of documents, is an important task that can help people to find information on these huge resources. Text categorization presents unique challenges due to the large number of attributes present in the data set, large number of training samples, attribute dependency, and multi-modality of categories. This has led to the development of a variety of text categorization algorithms [9, 10, 1, 22] that address these challenges to varying degrees.

In this paper we present a simple centroid-based document classification algorithm that is a special instance of Rocchio relevance feedback method [17]. In this algorithm, a centroid vector is computed to represent the documents of each class, and a new document is assigned to the class that corresponds to its most similar centroid vector, as measured by the cosine function. Extensive experiments presented in Section 3 show that this centroid-based classifier consistently and substantially outperforms other algorithms such as Naive Bayesian [13], k -nearest-neighbors [22], and C4.5 [16], on a wide range of datasets.

The primary focus of this paper is to explain this robust performance by analyzing the classification model used in this algorithm in contrasting it against those used in another algorithms. Our analysis shows that the similarity measure used by the centroid-based scheme allows it to classify a new document based on how closely its behavior matches the behavior of the documents belonging to different classes, as measured by the average similarity between the documents. This matching allows it to dynamically adjust for classes with different densities. Our analysis also shows that the similarity measure of the centroid-based scheme can account for dependencies between the terms in the different classes. We believe that this feature of the centroid-based classifier is the reason why it consistently outperforms the Naive Bayesian classifier, which can not take these dependencies into account.

The remainder of the paper is organized as follows. Section 3 experimentally evaluates this algorithm on a variety of data sets. Section 4 analyzes the classification model of the centroid-based classifier and compares it against those used by other algorithms. Finally, Section 5 provides directions for future research.

2 Centroid-Based Document Classifier

In the centroid-based classification algorithm, the documents are represented using the vector-space model [18]. In this model, each document d is considered to be a vector in the term-space. In its simplest form, each document is represented by the *term-frequency* (TF) vector $\mathbf{d}_{tf} = (tf_1, tf_2, \dots, tf_n)$, where tf_i is the frequency of the i th term in the document. A widely used refinement to this model is to weight each term based on its *inverse document frequency* (IDF) in the document collection. The motivation behind this weighting is that terms appearing frequently in many documents have limited

discrimination power, and for this reason they need to be de-emphasized. This is commonly done [18] by multiplying the frequency of each term i by $\log(N/df_i)$, where N is the total number of documents in the collection, and df_i is the number of documents that contain the i th term (i.e., document frequency). This leads to the *tf-idf* representation of the document, i.e., $\mathbf{d}_{tfidf} = (tf_1 \log(N/df_1), tf_2 \log(N/df_2), \dots, tf_n \log(N/df_n))$. Finally, in order to account for documents of different lengths, the length of each document vector is normalized so that it is of unit length, i.e., $\|\mathbf{d}_{tfidf}\|_2 = 1$. In the rest of the paper, we will assume that the vector representation \mathbf{d} of each document d has been weighted using *tf-idf* and it has been normalized so that it is of unit length.

In the vector-space model, the similarity between two documents d_i and d_j is commonly measured using the cosine function [18], given by

$$\cos(\mathbf{d}_i, \mathbf{d}_j) = \frac{\mathbf{d}_i \cdot \mathbf{d}_j}{\|\mathbf{d}_i\|_2 * \|\mathbf{d}_j\|_2}, \quad (1)$$

where “ \cdot ” denotes the dot-product of the two vectors. Since the document vectors are of unit length, the above formula simplifies to $\cos(\mathbf{d}_i, \mathbf{d}_j) = \mathbf{d}_i \cdot \mathbf{d}_j$.

Given a set S of documents and their corresponding vector representations, we define the **centroid** vector \mathbf{C} to be

$$\mathbf{C} = \frac{1}{|S|} \sum_{d \in S} \mathbf{d}, \quad (2)$$

which is nothing more than the vector obtained by averaging the weights of the various terms present in the documents of S . We will refer to the S as the **supporting set** for the centroid \mathbf{C} . Analogously to documents, the similarity between two centroid vectors and between a document and a centroid vector are computed using the cosine measure. In the first case,

$$\cos(\mathbf{C}_i, \mathbf{C}_j) = \frac{\mathbf{C}_i \cdot \mathbf{C}_j}{\|\mathbf{C}_i\|_2 * \|\mathbf{C}_j\|_2}, \quad (3)$$

whereas in the second case,

$$\cos(\mathbf{d}, \mathbf{C}) = \frac{\mathbf{d} \cdot \mathbf{C}}{\|\mathbf{d}\|_2 * \|\mathbf{C}\|_2} = \frac{\mathbf{d} \cdot \mathbf{C}}{\|\mathbf{C}\|_2}. \quad (4)$$

Note that even though the document vectors are of length one, the centroid vectors will not necessarily be of unit length.

The idea behind the centroid-based classification algorithm is extremely simple. For each set of documents belonging to the same class, we compute their centroid vectors. If there are k classes in the training set, this leads to k centroid vectors $\{\mathbf{C}_1, \mathbf{C}_2, \dots, \mathbf{C}_k\}$, where each \mathbf{C}_i is the centroid for the i th class. The class of a new document x is determined as follows. First we use the document-frequencies of the various terms computed from the training set to compute the *tf-idf* weighted vector-space representation of x , and scale it so x is of unit length. Then, we compute the similarity between x to all k centroids using the cosine measure. Finally, based on these similarities, we assign x to the class corresponding to the most similar centroid. That is, the class of x is given by

$$\arg \max_{j=1, \dots, k} (\cos(x, \mathbf{C}_j)). \quad (5)$$

The computational complexity of the learning phase of this centroid-based classifier is linear on the number of documents and the number of terms in the training set. The computation of the vector-space representation of the documents can be easily computed by performing at most three passes through the training set. Similarly, all k centroids can be computed in a single pass through the training set, as each centroid is computed by averaging the documents of the corresponding class. Moreover, the amount of time required to classify a new document x is at most $O(km)$, where m is the number of terms present in x . Thus, the overall computational complexity of this algorithm is very low, and is identical to fast document classifiers such as Naive Bayesian.

3 Experimental Results

We evaluated the performance of the centroid-based classifier by comparing against the naive Bayesian, C4.5, and k -nearest-neighbor classifiers on a variety of document collections. We obtained the naive Bayesian results using the Rainbow [14] with the multinomial event model [13]. The C4.5 results were obtained using a locally modified version of the C4.5 algorithm capable of handling sparse data sets. Finally, the k -nearest-neighbor results were obtained by using the *tf-idf* vector-space representation of the documents (identical to that used by the centroid-based classification algorithm), and using the number of neighbors $k = 10$.

3.1 Document Collections

The detailed characteristics of the various document collections used in our experiments are available in [7]¹. Note that for all data sets, we used a stop-list to remove common words, and the words were stemmed using Porter’s suffix-stripping algorithm [15]. Furthermore, we selected documents such that each document has only one class (or label). In other words, given a set of classes, we collected documents that have only one class from the set.

The first three data sets *west1*, *west2*, *west3* are from the statutory collections of the legal document publishing division of West Group described in [5]. Data sets *tr11*, *tr12*, *tr21*, *tr23*, *tr31*, *tr41*, *tr45*, *fbis*, *la1*, *la2*, *la12*, and *new3* are derived from TREC-5 [19], TREC-6 [19], and TREC-7 [19] collections. Data sets *re0* and *re1* are from Reuters-21578 text categorization test collection Distribution 1.0 [12]. We removed dominant classes such as “earn” and “acq” that have been shown to be relatively easy to classify. We then divided the remaining classes into 2 sets. Data sets *oh0*, *oh5*, *oh10*, *oh15*, and *ohscal* are from OHSUMED collection [8] subset of MEDLINE database. Data set *wap* is from the WebACE project (WAP) [2]. Each document corresponds to a web page listed in the subject hierarchy of Yahoo! [21].

¹ These data sets are available from <http://www.cs.umn.edu/~han/data/tmdata.tar.gz>.

3.2 Classification Performance

Since we constructed document data sets such that each document has single class label, we were able to perform true classification experiments where each document is classified to be any one of the classes and the performance is measured using classification accuracy. Note that the experimental results reported here are not directly comparable to other experiments reported in [4, 1, 9, 10, 22], because other experiments are based on binary classification per class and the performance is measured using precision and recall.

The classification accuracy of the various algorithms on the different data sets in our experimental testbed are shown in Table 1. These results correspond to the average classification accuracies of 10 experiments. In each experiment 80% of the documents were randomly selected as the training set, and the remaining 20% as the test set. The first three rows of this table show the results for the naive Bayesian, C4.5, and k -nearest neighbor schemes, whereas the last row shows the results achieved by the centroid-based classification algorithm (denoted as “Cntr” in the table). For each one of the data sets, we used a boldface font to highlight the algorithm that achieved the highest classification accuracy.

	west1	west2	west3	oh0	oh5	oh10	oh15	re0	re1	tr11	tr12	tr21	tr23	tr31	tr41	tr45	la1	la2	la12	fbis	wap	ohscal	new3
NB	86.7	76.5	75.1	89.1	87.1	81.2	84.0	81.1	80.5	85.3	79.8	59.6	69.3	94.1	94.5	84.7	87.6	89.9	89.2	77.9	80.6	74.6	74.4
C4.5	85.5	75.3	73.5	82.8	79.6	73.1	75.2	75.8	77.9	78.2	79.2	81.3	90.7	93.3	89.6	91.3	75.2	77.3	79.4	73.6	68.1	71.5	73.5
k NN	82.9	77.2	76.1	84.4	85.6	77.5	81.7	77.9	78.9	85.3	85.7	89.2	81.7	93.9	93.5	91.1	82.7	84.1	85.2	78.0	75.1	62.5	67.9
Cntr	87.5	79.0	81.6	89.3	88.2	85.3	87.4	79.8	80.4	88.2	90.3	91.6	85.2	94.9	95.7	92.9	87.4	88.4	89.1	80.1	81.3	75.4	79.7

Table 1. The classification accuracy achieved by the different classification algorithms.

Looking at the results of Table 1, we can see that naive Bayesian outperforms the other schemes in five out of the 23 data sets, C4.5 does better in one, the centroid-based scheme does better in 17, whereas the k -nearest-neighbor algorithm never outperforms the other schemes.

A more accurate comparison of the different schemes can be obtained by looking at what extent the performance of a particular scheme is statistically different from that of another scheme. We used two different statistical tests to compare the accuracy results obtained by the different classifiers. The first test is based on the resampled paired t test, and the second test is based on the sign test. A brief description of these tests is described in [7].

The statistical significance results using the resampled paired t test are summarized in Table 2, in which for each pair of classification algorithms, it shows the number of data sets that one performs statistically better, worse, or similarly than the other. Looking at this table, we can see that the centroid-based scheme compared to naive Bayesian, does better in ten data sets, worse in one data set, and they are statistically similar in twelve data sets. Similarly, compared to k NN, it does better in twenty, and it is statistically similar in three data sets. Finally, compared to C4.5, the centroid-based scheme does better in eighteen, worse in one, and statistically similar in four data sets.

	NB	k NN	C4.5
Cntr	10/1/12	20/0/3	18/1/4
NB		12/4/7	15/3/5
k NN			13/3/7

Table 2. Statistical comparison of different classification algorithms using the resampled paired t test. The entries in the table show the number of data sets that the classifier in the row performs better, worse or similarly than the classifier in the column.

The statistical significance results using the sign test are summarized in Table 3, in which for each pair of classification algorithms, it shows the z value. The z value was computed based on the average classification accuracy of 10 trials. A z value greater than 1.96, indicates that the classifier of the row is statistically better than the classifier of the column. Looking at this table, we can see that the centroid-based scheme does better than naive Bayesian, k NN, and C4.5. Naive Bayesian does better than C4.5, but does similarly with respect to k NN. Finally, k NN does better than C4.5.

From these results, we can see that the simple centroid-based classification algorithm outperforms all remaining schemes, with naive Bayesian being second, k -nearest-neighbor being third, and C4.5 being the last. Note that the better performance of NB and k NN over decision tree classification algorithms such as C4.5 agrees to results reported in [3, 22] using precision and recall of binary classification.

Recently, Support Vector Machines (SVM) has been shown to be very effective in text classification [9]. We were not able to directly compare the centroid-based scheme with the SVM, because the SVM code used in [9] was written for binary classification only. We plan to perform comparison studies between SVM and the centroid-based scheme by performing binary classification in the future.

4 Analysis

4.1 Classification Model

The surprisingly good performance of the centroid-based classification scheme suggests that it employs a sound underlying classification model. The goal of this section is to understand this classification model and compare it against those used by other schemes.

	NB	k NN	C4.5
Cntr	2.71	4.80	4.38
NB		1.46	3.54
k NN			2.71

Table 3. Statistical comparison of different classification algorithms using the sign test. The values in the table are z values and value greater than 1.96 shows that the classifier of the row is statistically better than the classifier of the column.

In order to understand this model we need to understand the formula used to determine the similarity between a document x , and the centroid vector C of a particular class (Equation 4), as this computation is essential in determining the class of x (Equation 5). From Equation 4, we see that the similarity (i.e., cosine) between x and C is the ratio of the dot-product between x and C divided by the length of C . If S is the set of documents used to create C , then from Equation 2, we have that:

$$x \cdot C = x \cdot \left(\frac{1}{|S|} \sum_{d \in S} d \right) = \frac{1}{|S|} \sum_{d \in S} x \cdot d = \frac{1}{|S|} \sum_{d \in S} \cos(x, d).$$

That is, the dot-product is the average similarity (as measured by the cosine function) between the new document x and all other documents in the set. The meaning of the length of the centroid vector can also be easily understood using the fact that $\|C\|_2 = \sqrt{C \cdot C}$. Then, from Equation 2 we have that:

$$\|C\|_2 = \sqrt{\left(\frac{1}{|S|} \sum_{d \in S} d \right) \cdot \left(\frac{1}{|S|} \sum_{d \in S} d \right)} = \sqrt{\frac{1}{|S|^2} \sum_{d_i \in S} \sum_{d_j \in S} \cos(d_i, d_j)}.$$

Hence, the length of the centroid vector is the square-root of the average pairwise similarity between the documents that support the centroid. There are two things to be noted about this formula; first, this average similarity also includes the self-similarity between the documents in the supporting set; second, because all the documents have been scaled to be of unit length, the length of the centroid vector will always be less or equal to one. In summary, the similarity between a test document and the centroid vector of a particular class, is nothing more than the average similarity between the test document and all the documents in that class, divided by the square-root of the average similarity between the documents in the class itself.

The above discussion provides us with a qualitative understanding on how the centroid scheme determines the similarity between a test document and a particular class. Essentially, it computes the average similarity between the test document and all the other documents in that class, and then it amplifies that similarity, based on how similar to each other are the documents of that class. If the average pairwise similarity between the documents of the class is small (i.e., the class is *loose*), then that amplification is higher, whereas if the average pairwise similarity is high (i.e., the class is *tight*), then this amplification is smaller.

To better understand this classification model consider the following simple binary classification algorithm, that we will refer to it as \mathcal{H} . Let A and B be the two classes, let \bar{S}_A be the average similarity between the items in A , \bar{S}_B be the average similarity between the items in B , and let $\bar{S}_{A,B}$ be the average similarity between all the items (a, b) such that $a \in A$, and $b \in B$. Now consider a test item x , and let $\bar{S}_{x,A}$, and $\bar{S}_{x,B}$ be the average similarities between x and all the items in A and B , respectively. This setting is illustrated in Figure 1. In this classifier, x will be classified as either A or B based on how closely its behavior matches the behavior of the items in class A and the items in class B , as measured by their average similarities.

This behavior can be modeled by looking at the ratios $\bar{S}_A / \bar{S}_{A,B}$ and $\bar{S}_B / \bar{S}_{A,B}$, and comparing them against the ratios $\bar{S}_{x,A} / \bar{S}_{x,B}$ and $\bar{S}_{x,B} / \bar{S}_{x,A}$. The first of these ratios




Fig. 1. A simple binary classifier.

$(\bar{S}_A/\bar{S}_{A,B})$ measures how much *stronger* is the internal similarity between items belonging to class A relative to their similarity to items belonging to class B . Similarly, the second ratio $(\bar{S}_B/\bar{S}_{A,B})$ measures how much stronger is the internal similarity between items belonging to class B relative to their similarity to items belonging to class A . Finally, the last two ratios, measure how much stronger is the similarity of x to the items in A compared to the items in B , and vice-versa. Given the above ratios, then the classification algorithm \mathcal{H} will assign x to class A iff,

$$\frac{\bar{S}_{x,A}/\bar{S}_{x,B}}{\bar{S}_A/\bar{S}_{A,B}} \geq \frac{\bar{S}_{x,B}/\bar{S}_{x,A}}{\bar{S}_B/\bar{S}_{A,B}}, \quad (6)$$

otherwise it will assign in to class B . Essentially, \mathcal{H} compares the strength of the similarity of x to class A relative to the strength of the similarity of items already in A (left side of the inequality), against the strength of the similarity of x to class B relative to the strength of the similarity of items already in B (right side of the inequality), and assigns x to the class for which the relative strength is higher. Performing some simple algebraic manipulations in Equation 6, and canceling out the $\bar{S}_{A,B}$ terms that appear on both side of the inequality we have that:

$$\frac{\bar{S}_{x,A}/\bar{S}_{x,B}}{\bar{S}_A/\bar{S}_{A,B}} \geq \frac{\bar{S}_{x,B}/\bar{S}_{x,A}}{\bar{S}_B/\bar{S}_{A,B}} \Rightarrow \frac{\bar{S}_{x,A}^2}{\bar{S}_A} \geq \frac{\bar{S}_{x,B}^2}{\bar{S}_B} \Rightarrow \frac{\bar{S}_{x,A}}{\sqrt{\bar{S}_A}} \geq \frac{\bar{S}_{x,B}}{\sqrt{\bar{S}_B}}. \quad (7)$$

We can extend \mathcal{H} to problems with more than two classes, by using a *tournament* method, and thus assigning x to the class for which $\bar{S}_{x,j}/\sqrt{\bar{S}_j}$ is the highest among all classes j .

Now, from the earlier discussion, we know that in the case in which the data items in the above problem are unit-length document vectors, and the similarity is computed using the cosine measure, then from Equation 7 we have that \mathcal{H} will assign x to class A , iff

$$\cos(x, \mathbf{C}_A) \geq \cos(x, \mathbf{C}_B),$$

otherwise x will be assigned to class B ; where \mathbf{C}_A and \mathbf{C}_B are the centroid vectors of class A and B , respectively. Thus, the classification model used by the centroid-based document classifier is identical to that used by \mathcal{H} , that is, it assigns a new document x to the class whose documents better match the behavior of x , as measured by average document similarities.

4.2 Comparison With Other Classifiers

One of the advantages of the centroid-based scheme is that it summarizes the characteristics of each class, in the form of the centroid vector. A similar summarization is also performed by naive Bayesian, in the form of the per-class term-probability distribution functions.

The advantage of the summarization performed by the centroid vectors is that it combines multiple prevalent features together, even if these features are not simultaneously present in a single document. That is, if we look at the prominent dimensions of the centroid vector (i.e., highest weight terms), these will correspond to terms that appear frequently in the documents of the class, but not necessarily all in the same set of documents. This is particularly important for high dimensional data sets for which the coverage of any individual feature is often quite low. Moreover, in the case of documents, this summarization has the additional benefit of addressing issues related to synonyms, as commonly used synonyms will be represented in the centroid vector (see [7] for some of the centroid vectors of data sets used in the experiments). For these reasons, the centroid-based classification algorithm (as well as naive Bayesian) tend to perform better than the C4.5 and the k -nearest neighbor classification algorithms.

The better performance of the centroid-based scheme over the naive Bayesian classifier is due to the method used to compute the similarity between a test document and a class. In the case of naive Bayesian, this is done using Bayes rule, assuming that when conditioned on each class, the occurrence of the different terms is independent. However, this is far from being true in real document collections [11]. One way of understanding the dependence between terms is to look at the degree at which various terms co-occur in the documents of a particular class. If the degree of term co-occurrence is high, then these terms are positively dependent, as the probability of seeing one of the co-occurring terms is high provided that we have seen one of the other co-occurring terms. As the degree of term co-occurrence decreases, the positive dependence also decreases, and after a certain point it gives rise to negative dependence among the terms. In this case, the conditional probability of seeing a certain term is high provided that we have not seen some other terms. The existence of such positive and negative dependence between terms of a particular class causes naive Bayesian to compute a distorted estimate of the probability that a particular document belongs to that class. If there is positive dependence between the terms in the class, then the probability estimate will be higher than it actually is, whereas if there is negative dependence between the terms, then the probability estimate will be smaller than it actually is. Unfortunately, naive Bayesian has no way by which to account for such term dependence, and much more complicated classifiers such as Bayesian Networks need to be used [6].

On the other hand, the similarity function used by the centroid-based scheme does account for term dependence within each class. From the discussion in Section 4, we know that the similarity of a new document x to a particular class is computed as the ratio of two quantities. The first is the average similarity of x to all the documents in the class, and the second is the square-root of the average similarity of the documents within the class. To a large extent, the first quantity is very similar, in character, to the probability estimate used by the naive Bayesian algorithm, and it suffers from similar over- and under-estimation problems in the case of term dependence. As in the case

of naive Bayesian, if the class contains terms that are positively dependent, then the average similarity of x to the documents in the class will be high, as it will tend to match most of the co-occurring terms. Similarly, if the class contains negatively dependent terms, then the average similarity of x to the documents in the class will be small as it will be unnecessarily penalized for not matching the negatively dependent terms.

However, the second quantity of the similarity function, (i.e., the square-root of the average similarity of the documents within the class) does account for term dependency. This average similarity depends on the degree at which terms co-occur in the different documents. In general, if the average similarity between the documents of a class is high, then the documents have a high degree of term co-occurrence (since the similarity between a pair of documents computed by the cosine function, is high when the documents have similar set of terms). On the other hand, as the average similarity between the documents decreases, the degree of term co-occurrence also decreases. Since this average internal similarity is used to amplify the similarity between a test document and the class, this amplification is minimal when there is a large degree of positive dependence among the terms in the class, and increases as the positive dependence decreases. Consequently, this amplification acts as a correction parameter to account for the over- and under-estimation of the similarity that is computed by the first quantity in the document-to-centroid similarity function. We believe that this feature of the centroid-based classification scheme is the reason that it outperforms the naive Bayesian classifier in the experiments shown in Section 3.

5 Discussion & Concluding Remarks

In this paper we focused on a simple linear-time centroid-based document classification algorithm. Our experimental evaluation has shown that the centroid-based classifier consistently and substantially outperforms other classifiers on a wide range of data sets. We have shown that the power of this classifier is due to the function that it uses to compute the similarity between a test document and the centroid vector of the class. This similarity function can account for both the term similarity between the test document and the documents in the class, as well as for the dependencies between the terms present in these documents.

There are many ways to further improve the performance of this centroid-based classification algorithm. First, in its current form it is not well suited to handle multi-modal classes. However, support for multi-modality can be easily incorporated by using a clustering algorithm to partition the documents of each class into multiple subsets, each potentially corresponding to a different mode, or using similar techniques to those used by the generalized instance set classifier [10]. Second, the classification performance can be further improved by using techniques that adjust the importance of the different features in a supervised setting. A variety of such techniques have been developed in the context of k -nearest-neighbor classification [20], all of which can be extended to the centroid-based classifier.

References

1. L. Baker and A. McCallum. Distributional clustering of words for text classification. In *SIGIR-98*, 1998.
2. D. Boley, M. Gini, R. Gross, E.H. Han, K. Hastings, G. Karypis, V. Kumar, B. Mobasher, and J. Moore. Document categorization and query generation on the world wide web using WebACE. *AI Review (accepted for publication)*, 1999.
3. W.W. Cohen. Fast effective rule induction. In *Proc. of the Twelfth International Conference on Machine Learning*, 1995.
4. W.W. Cohen and H. Hirsh. Joins that generalize: Text classification using WHIRL. In *Proc. of the Fourth Int'l Conference on Knowledge Discovery and Data Mining*, 1998.
5. T. Curran and P. Thompson. Automatic categorization of statute documents. In *Proc. of the 8th ASIS SIG/CR Classification Research Workshop*, Tucson, Arizona, 1997.
6. N. Friedman, D. Geiger, and M. Goldszmidt. Bayesian network classifiers. *Machine Learning*, 29:131–163, 1997.
7. E.H. Han and G. Karypis. Centroid-based document classification algorithms: Analysis & experimental results. Technical Report TR-00-017, Department of Computer Science, University of Minnesota, Minneapolis, 2000. Available on the WWW at URL <http://www.cs.umn.edu/~karypis>.
8. W. Hersh, C. Buckley, T.J. Leone, and D. Hickam. OHSUMED: An interactive retrieval evaluation and new large test collection for research. In *SIGIR-94*, pages 192–201, 1994.
9. T. Joachims. Text categorization with support vector machines: Learning with many relevant features. In *Proc. of the European Conference on Machine Learning*, 1998.
10. Wai Lam and Chao Yang Ho. Using a generalized instance set for automatic text categorization. In *SIGIR-98*, 1998.
11. D. Lewis. Naive (bayes) at forty: The independence assumption in information retrieval. In *Tenth European Conference on Machine Learning*, 1998.
12. D. D. Lewis. Reuters-21578 text categorization test collection distribution 1.0. <http://www.research.att.com/~lewis>, 1999.
13. A. McCallum and K. Nigam. A comparison of event models for naive bayes text classification. In *AAAI-98 Workshop on Learning for Text Categorization*, 1998.
14. Andrew Kachites McCallum. Bow: A toolkit for statistical language modeling, text retrieval, classification and clustering. <http://www.cs.cmu.edu/mccallum/bow>, 1996.
15. M. F. Porter. An algorithm for suffix stripping. *Program*, 14(3):130–137, 1980.
16. J. Ross Quinlan. *C4.5: Programs for Machine Learning*. Morgan Kaufmann, San Mateo, CA, 1993.
17. J.J. Jr. Rocchio. The SMART retrieval system: Experiments in automatic document processing. In Gerard Salton, editor, *Relevance feedback in information retrieval*. Prentice-Hall, Inc., 1971.
18. G. Salton. *Automatic Text Processing: The Transformation, Analysis, and Retrieval of Information by Computer*. Addison-Wesley, 1989.
19. TREC. Text REtrieval conference. <http://trec.nist.gov>.
20. D. Wettschereck, D.W. Aha, and T. Mohri. A review and empirical evaluation of feature-weighting methods for a class of lazy learning algorithms. *AI Review*, 11, 1997.
21. Yahoo! Yahoo! <http://www.yahoo.com>.
22. Y. Yang and X. Liu. A re-examination of text categorization methods. In *SIGIR-99*, 1999.